

Pro_Training

February 24, 2024

0.1 LOAD LIBRARIES

```
[17]: import numpy as np
import tensorflow as tf
from tensorflow.keras import models, layers
import matplotlib.pyplot as plt
import os
```

```
[18]: Image_Size = 256
BATCH_SIZE = 32
Channels = 3
EPOCHS = 20
```

0.2 This Code Loads the Data Into Tensorflow Database

```
[19]: Dataset = tf.keras.preprocessing.image_dataset_from_directory(
    "Data",
    shuffle = True,
    image_size=(Image_Size,Image_Size),
    batch_size = BATCH_SIZE
)
```

Found 248 files belonging to 3 classes.

0.3 Analyse the DATA

```
[20]: Class_names = Dataset.class_names
```

```
[21]: Class_names
```

```
[21]: ['Blackpod', 'Frosty', 'Healthy']
```

```
[22]: len(Dataset)
```

```
[22]: 8
```

```
[23]: 9*32
```

[23]: 288

```
[24]: plt.figure(figsize=(10,10))
for image_batch, label_batch in Dataset.take(1):
    for i in range(12):
        ax = plt.subplot(3,4,i+1)

        #print(image_batch.shape)
        #print(label_batch.numpy())
        #print(image_batch[0].numpy)## Changing tensor to a numpy
        #print(image_batch[0].shape)
        plt.imshow(image_batch[i].numpy().astype("uint8"))
        plt.axis("off")
        plt.title(Class_names[label_batch[i]])
```

Blackpod



Frosty



Healthy



Healthy



Healthy



Blackpod



Blackpod



Blackpod



Healthy



Healthy



Healthy



Healthy



```
[25]: len(Dataset)
```

```
[25]: 8
```

0.4 Splitting the DATASET

```
[26]: #80% ==> training
      #20% ==> 10% validation, 10% test to measure the accuracy of the model
```

0.4.1 Using Dataset.take to Split the DATA

```
[27]: train_size = 0.8
      len(Dataset)* train_size # getting the percentage of the train data size from
      ↪the whole data
```

```
[27]: 6.4
```

```
[28]: train_ds = Dataset.take(7)
      len(train_ds)## Taking the train size percentage from the Data
```

```
[28]: 7
```

```
[29]: test_ds = Dataset.skip(7)
      len(test_ds)## Skipping the train size to get the test size
```

```
[29]: 1
```

```
[30]: val_size =0.1
      len(Dataset)*val_size## splitting the test size into Validation dataset and
      ↪test dataset
```

```
[30]: 0.8
```

```
[31]: val_ds= test_ds.take(1)
      len(val_ds)
```

```
[31]: 1
```

```
[32]: test_ds = test_ds.skip(1)
      len(test_ds)
```

```
[32]: 0
```

0.5 Splitting the DATASET

```
[33]: def get_dataset_partitions_tf(ds, train_split = 0.8, val_split =0.1, test_split=
      ↪0.1, shuffle =True, shuffle_size =10000):
      ds_size = len(ds)
      if shuffle:
          ds = ds.shuffle(shuffle_size, seed = 12)
      train_size =int(train_split *ds_size)
      val_size = int(val_split * ds_size)

      train_ds = ds.take(train_size)## Taking the train size from the dataset
      val_ds = ds.skip(train_size).take(val_size)# Skip the train_size and the
      ↪remaining 20% take Val_size

      test_ds = val_ds= ds.skip(train_size).skip(val_size)## skip both train and
      ↪vals_size and the remaining is Test_size
      return train_ds, val_ds, test_ds
```

```
[34]: train_ds, val_ds, test_ds = get_dataset_partitions_tf(Dataset)
```

```
[35]: len(train_ds)
```

```
[35]: 6
```

```
[36]: len(val_ds)
```

```
[36]: 2
```

```
[37]: len(test_ds)
```

```
[37]: 2
```

0.6 Caching to improve the performance of the pipeline

Shuffle 1000 will shuffle the images

Prefetch to loads the next set of batch from the disk to improve performance

```
[38]: train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size =tf.data.
      ↪AUTOTUNE)

      val_ds = val_ds.cache().shuffle(1000).prefetch(buffer_size =tf.data.AUTOTUNE)

      test_ds = test_ds.cache().shuffle(1000).prefetch(buffer_size =tf.data.AUTOTUNE)
      ↪## Necessary for training performance
```

0.6.1 Preprocessing Resizing and Rescaling

```
[39]: resize_and_rescale = tf.keras.Sequential([
        layers.experimental.preprocessing.Rescaling(Image_Size,Image_Size),
        layers.experimental.preprocessing.Rescaling(1.0/255) ## Rescaling the
        ↪ images to 255
    ])
```

0.7 Creating more samples due to the fewer images to maximize the variables for effective prediction

```
[40]: data_augmentaion = tf.keras.Sequential([
        layers.experimental.preprocessing.RandomFlip("horizontal_and_vertical"),
        layers.experimental.preprocessing.RandomRotation(0.2), ## Randomflip and
        ↪ some rotation to have a diverse forms of the data
        layers.experimental.preprocessing.RandomZoom(0.1), ## to make the model a
        ↪ robust one.
        layers.experimental.preprocessing.RandomContrast(0.1)
    ])
```

0.8 Build First Classifier (CNN)

```
[41]: input_shape = (BATCH_SIZE,Image_Size,Image_Size,Channels)
n_classes = 3

model = models.Sequential([
    resize_and_rescale,
    data_augmentaion,
    layers.Conv2D(32,(3,3), activation = 'relu',input_shape =input_shape), ##
    ↪ Need to have a lot of layers in order for the prediction to be inact
    layers.MaxPooling2D((2,2)), ## this helps to scans over the image to pull
    ↪ out the max values of the image

    layers.Conv2D(64, kernel_size = (3,3), activation = 'relu'),
    layers.MaxPooling2D((2,2)),

    layers.Conv2D(64, kernel_size = (3,3), activation = 'relu'),
    layers.MaxPooling2D((2,2)),

    layers.Conv2D(64,(3,3), activation = 'relu'),
    layers.MaxPooling2D((2,2)),

    layers.Conv2D(64,(3,3), activation = 'relu'),
    layers.MaxPooling2D((2,2)),
    layers.Flatten(), ## flatten
    layers.Dense(64, activation = 'relu'), # and add a densed layer
])
```

```

        layers.Dense(n_classes, activation = 'softmax'))])# softmax activation
        ↪function normalizes the probability of the classes

model.build(input_shape = input_shape) ## Force defining the Nero achitecture

```

```
[42]: model.summary()### Module achitecture
```

Model: "sequential_2"

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| ===== | | |
| sequential (Sequential) | (32, 256, 256, 3) | 0 |
| sequential_1 (Sequential) | (32, 256, 256, 3) | 0 |
| conv2d (Conv2D) | (32, 254, 254, 32) | 896 |
| max_pooling2d (MaxPooling2D) | (32, 127, 127, 32) | 0 |
| conv2d_1 (Conv2D) | (32, 125, 125, 64) | 18496 |
| max_pooling2d_1 (MaxPooling2D) | (32, 62, 62, 64) | 0 |
| conv2d_2 (Conv2D) | (32, 60, 60, 64) | 36928 |
| max_pooling2d_2 (MaxPooling2D) | (32, 30, 30, 64) | 0 |
| conv2d_3 (Conv2D) | (32, 28, 28, 64) | 36928 |
| max_pooling2d_3 (MaxPooling2D) | (32, 14, 14, 64) | 0 |
| conv2d_4 (Conv2D) | (32, 12, 12, 64) | 36928 |
| max_pooling2d_4 (MaxPooling2D) | (32, 6, 6, 64) | 0 |
| flatten (Flatten) | (32, 2304) | 0 |
| dense (Dense) | (32, 64) | 147520 |
| dense_1 (Dense) | (32, 3) | 195 |
| ===== | | |

Total params: 277891 (1.06 MB)
Trainable params: 277891 (1.06 MB)
Non-trainable params: 0 (0.00 Byte)

0.9 Defining the optimizer, loss function and metrics

```
[43]: model.compile(  
    optimizer = 'adam',  
    loss = tf.keras.losses.SparseCategoricalCrossentropy(from_logits = False),  
    metrics = ['accuracy'])## accuracy is the metric used to track the training_  
    ↪process  
  
    #callback = keras.callbacks.EarlyStopping(monitor = 'val_loss',  
                                              # patience = 3,  
                                              # restore_best_weights = True)
```

0.10 Training the Network

```
[44]: history = model.fit(  
    train_ds,  
    epochs = EPOCHS,  
    batch_size = BATCH_SIZE,  
    verbose = 1,  
    validation_data = val_ds  
)
```

```
Epoch 1/20  
6/6 [=====] - 7s 875ms/step - loss: 13.2165 - accuracy:  
0.3913 - val_loss: 1.4085 - val_accuracy: 0.2812  
Epoch 2/20  
6/6 [=====] - 5s 852ms/step - loss: 1.2479 - accuracy:  
0.3315 - val_loss: 1.0224 - val_accuracy: 0.4688  
Epoch 3/20  
6/6 [=====] - 5s 875ms/step - loss: 0.8919 - accuracy:  
0.5761 - val_loss: 0.9722 - val_accuracy: 0.5781  
Epoch 4/20  
6/6 [=====] - 6s 978ms/step - loss: 0.8996 - accuracy:  
0.5761 - val_loss: 0.9307 - val_accuracy: 0.5938  
Epoch 5/20  
6/6 [=====] - 5s 915ms/step - loss: 0.8532 - accuracy:  
0.6087 - val_loss: 0.8799 - val_accuracy: 0.5781  
Epoch 6/20  
6/6 [=====] - 5s 903ms/step - loss: 0.7242 - accuracy:  
0.6957 - val_loss: 0.8324 - val_accuracy: 0.6562  
Epoch 7/20  
6/6 [=====] - 6s 963ms/step - loss: 0.7483 - accuracy:
```

```

0.6685 - val_loss: 0.7939 - val_accuracy: 0.6562
Epoch 8/20
6/6 [=====] - 5s 922ms/step - loss: 0.6945 - accuracy:
0.7011 - val_loss: 0.9994 - val_accuracy: 0.5469
Epoch 9/20
6/6 [=====] - 6s 1s/step - loss: 0.6763 - accuracy:
0.7011 - val_loss: 0.7517 - val_accuracy: 0.7031
Epoch 10/20
6/6 [=====] - 6s 993ms/step - loss: 0.6432 - accuracy:
0.6957 - val_loss: 0.8828 - val_accuracy: 0.6250
Epoch 11/20
6/6 [=====] - 6s 985ms/step - loss: 0.6038 - accuracy:
0.7391 - val_loss: 0.7323 - val_accuracy: 0.7344
Epoch 12/20
6/6 [=====] - 6s 1s/step - loss: 0.5418 - accuracy:
0.7554 - val_loss: 0.6752 - val_accuracy: 0.7812
Epoch 13/20
6/6 [=====] - 6s 976ms/step - loss: 0.5476 - accuracy:
0.7772 - val_loss: 1.0608 - val_accuracy: 0.5781
Epoch 14/20
6/6 [=====] - 6s 969ms/step - loss: 0.6713 - accuracy:
0.7011 - val_loss: 0.6956 - val_accuracy: 0.7500
Epoch 15/20
6/6 [=====] - 6s 920ms/step - loss: 0.6413 - accuracy:
0.6793 - val_loss: 1.1077 - val_accuracy: 0.5156
Epoch 16/20
6/6 [=====] - 6s 931ms/step - loss: 0.7253 - accuracy:
0.6359 - val_loss: 0.7474 - val_accuracy: 0.6719
Epoch 17/20
6/6 [=====] - 6s 952ms/step - loss: 0.5881 - accuracy:
0.7663 - val_loss: 0.6395 - val_accuracy: 0.7812
Epoch 18/20
6/6 [=====] - 6s 939ms/step - loss: 0.5265 - accuracy:
0.7717 - val_loss: 0.6395 - val_accuracy: 0.8281
Epoch 19/20
6/6 [=====] - 5s 887ms/step - loss: 0.4703 - accuracy:
0.8207 - val_loss: 0.8651 - val_accuracy: 0.6875
Epoch 20/20
6/6 [=====] - 5s 916ms/step - loss: 0.4997 - accuracy:
0.7989 - val_loss: 0.6610 - val_accuracy: 0.7969

```

0.10.1 To define how well the model is performed with a data that hasn't been seen by the model in order to avoid any bias

```

[45]: scores = model.evaluate(test_ds)## runing the model on the test_ds for the
      ↪first time (avoid bias)

```

```

2/2 [=====] - 1s 260ms/step - loss: 0.3474 - accuracy:

```


0.8281

```
[46]: scores
```

```
[46]: [0.3474399745464325, 0.828125]
```

```
[47]: history
```

```
[47]: <keras.src.callbacks.History at 0x1485a2ed0>
```

```
[48]: history.params
```

```
[48]: {'verbose': 1, 'epochs': 20, 'steps': 6}
```


```
[49]: history.history.keys()
```

```
[49]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

0.11 Plotting History

```
[50]: acc = history.history['accuracy']  
      val_acc = history.history['val_accuracy']  
  
      loss = history.history['loss']  
      val_loss = history.history['val_loss']
```

```
[51]: plt.figure(figsize=(10,10))  
      plt.subplot(1,2,1)  
      plt.plot(range(EPOCHS), acc, label = 'Training Accuracy')  
      plt.plot(range(EPOCHS), val_acc, label = 'Validation Accuracy')  
      plt.legend(loc='lower right')  
      plt.title('Training and Validation Accuracy') ## High accuracy was achieved  
  
      #plt.figure(figsize=(8,8))  
      plt.subplot(1,2,2)  
      plt.plot(range(EPOCHS), loss, label = 'Training Loss')  
      plt.plot(range(EPOCHS), val_loss, label = 'Validation Loss')  
      plt.legend(loc='upper right')  
      plt.title('Training and Validation Loss')  
      plt.show()
```



```
[52]: np.argmax([1.3001359e-05, 1.9586462e-04, 9.9979120e-01])
```

```
[52]: 2
```

0.12 Making a Prediction

```
[58]: plt.figure(figsize=(8,8))
for images_batch, labels_batch in test_ds.take(1):## taking one batch
    first_image = images_batch[4].numpy().astype('uint8')
    first_label = label_batch[4].numpy()

    print("First image to Predict")
```

```
plt.imshow(first_image)
print("Actual Label:", Class_names[first_label])

batch_prediction = model.predict(image_batch)
print("Predicted Label:", Class_names[np.argmax(batch_prediction[4])])
plt.axis('off')
```

First image to Predict

Actual Label: Healthy

1/1 [=====] - 0s 229ms/step

Predicted Label: Healthy

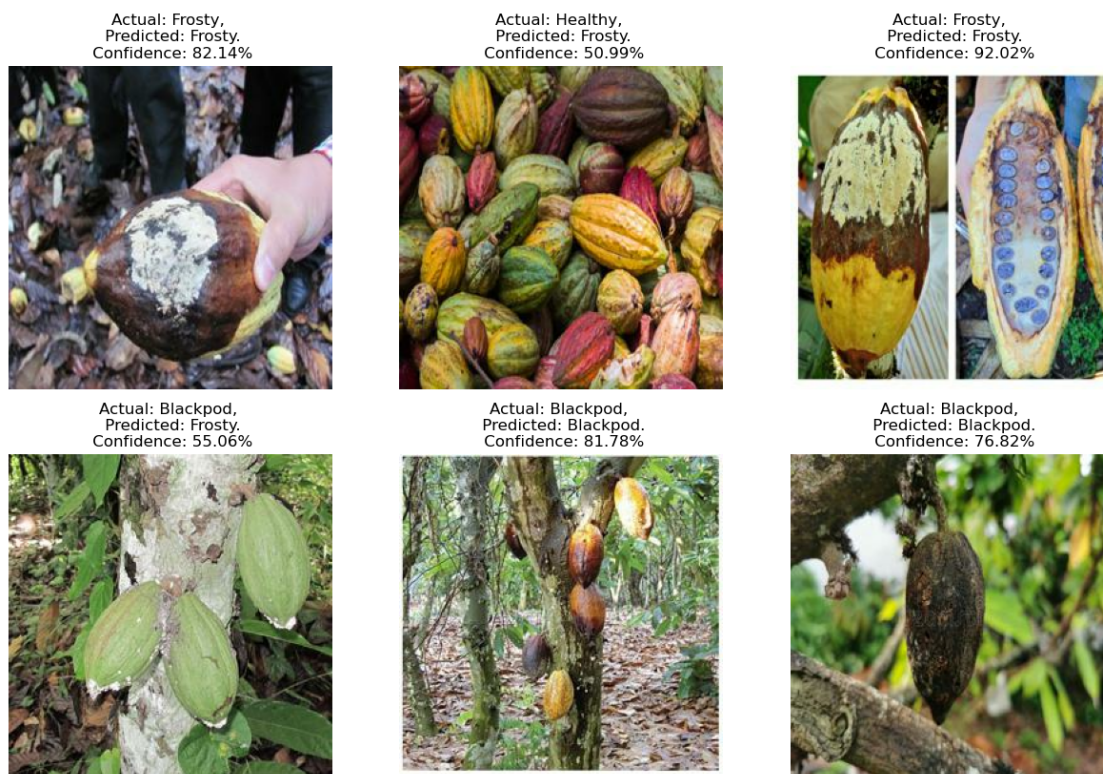


0.13 Function Determining the Predicted_Class/Confidence_Level of the Model

```
[54]: def predict(model,img):  
    img_array = tf.keras.preprocessing.image.img_to_array(images[i].numpy())  
    img_array = tf.expand_dims(img_array,0)# create batch  
  
    predictions = model.predict(img_array)  
  
    predicted_class = Class_names[np.argmax(predictions[0])]   
  
    Confidence = round(100* (np.max(predictions[0])),2)  
    return predicted_class, Confidence
```

```
[59]: plt.figure(figsize=(15,15))  
for images, labels in test_ds.take(1):## one batch  
    for i in range(6):## displaying only 9 images out if the batch  
        ax = plt.subplot(3,3,i+1)  
        plt.imshow(images[i].numpy().astype('uint8'))  
  
        Predicted_Class, Confidence = predict(model, images[i].numpy())  
        Actual_Class = Class_names[labels[i]]  
  
        plt.title(f"Actual: {Actual_Class}, \n Predicted: {Predicted_Class}.\n  
↪Confidence: {Confidence}%")  
  
        plt.axis('off')  
        plt.savefig('Fig')
```

```
1/1 [=====] - 0s 36ms/step  
1/1 [=====] - 0s 29ms/step  
1/1 [=====] - 0s 30ms/step  
1/1 [=====] - 0s 32ms/step  
1/1 [=====] - 0s 29ms/step  
1/1 [=====] - 0s 28ms/step
```



0.14 Saving the Model

```
[ ]: model_version = 1
model.save(f"../models/{model_version}")# model.. will take your from present_
↳ drectory to the new model directory
```

```
[ ]: #import os
#model_version = max([int(i) for i in os.listdir("../models")+ [0]])+1 #
↳ Changing a String to Integer
#model.save(f"../models/{model_version}")
```

```
[ ]:
```

```
[ ]:
```